

lab2 DPDK

Q1: What's the purpose of using hugepage?

网络传输是I/O密集型的时使用的连续空间很大，如果用正常的4K页来翻译，会需要很多个PTE，产生多次缺页中断，而使用更大的页，可以减少PTE数量，减少缺页中断的次数，并减少地址翻译过程的级数。缺页中断对性能的影响。同时TLB的命中率会提升很大。

Q2: Take examples/helloworld as an example, describe the execution flow of DPDK programs?

1. `rte_eal_init(argc,argv)` 将程序的参数传入init函数，对Environment Abstraction Layer(EAL)进行初始化
2. 通过下述方式让所有slaves执行lcore_hello函数。

```
RTE_LCORE_FOREACH_SLAVE(lcore_id) {  
    rte_eal_remote_launch(lcore_hello, NULL, lcore_id);  
}
```

3. 父节点自己执行lcore_hello函数
4. 通过 `rte_eal_mp_wait_lcore()` 这个函数等待其他结点结束执行

2,3两步可以简化为如下

```
rte_eal_mp_remote_launch(lcore_hello, NULL, CALL_MASTER);
```

Q3: Read the codes of examples/skeleton, describe DPDK APIs related to sending and receiving packets.

1. `rte_pktmbuf_pool_create` 创建mempool 用于保存message buffer(rte_mbuf)
2. port initialization
 - `rte_eth_rx_queue_setup` 函数创建receive queue
 - `rte_eth_tx_queue_setup()` 函数创建transmit queue
 - `rte_eth_dev_start(port)` 函数开启以太网端口
 - `rte_eth_promiscuous_enable(port)` 函数以混杂模式启动receive
3. 发送和接收

- `rte_eth_rx_burst` 接受一堆数据包到buf
- `rte_eth_tx_burst` 发送一堆数据包到buf

Q4: Describe the data structure of `rte_mbuf`

用来缓存数据包，通过`rte_pktmbuf_alloc`函数分配，保存在`rte_mempool`中。

`rte_eth_rx_burst` , `rte_eth_tx_burst` 都需要通过`rte_mbuf`结构来发送接收数据包。

利用`rte_pktmbuf_prepend`,`rte_pktmbuf_append`函数设置空间，并利用`rte_pktmbuf_mtod_offset`来获取指向特定offset的指针

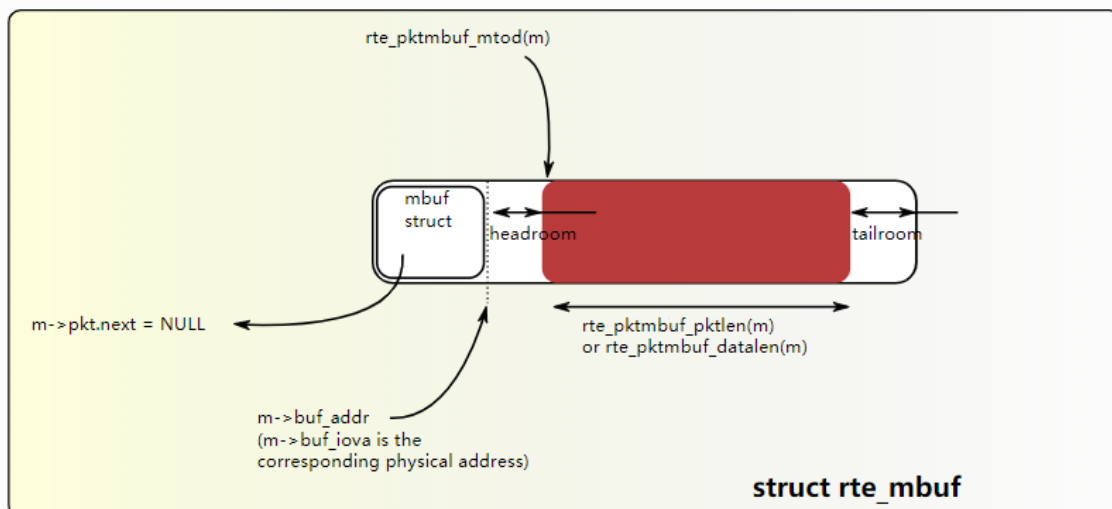


Fig. 9.1 An mbuf with One Segment

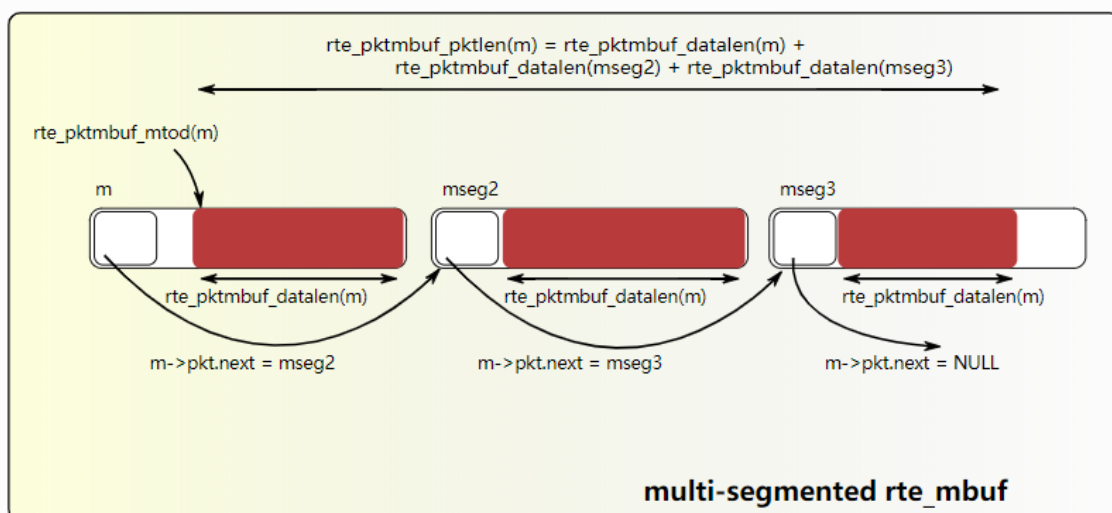


Fig. 9.2 An mbuf with Three Segments

Part2

编写

- 根据ether_hdr, ipv4_hdr, udp_hdr(似乎最新版本的dpdk加上前缀rte_)的结构逐一将对应包的信息填入。利用rte_pktmbuf_mtod_offset, rte_pktmbuf_mtod函数获取mbuf不同位置的指针, 并通过指针修改包的内容。填写时注意每一项都需要改成网络传输的big endian。注意各部分的length要计算正确。
- 在basicfwd的基础上, 修改其lcore_main函数。通过rte_pktmbuf_prepend预先设定长度, 并利用刚才写的函数构建udp包
- 利用rte_eth_tx_burst, 循环发送刚刚构建好的udp包

```
222         printf("Sending packet: %d\n", nb)
223         sleep(1);
224     }
225 }
226 /*
227  * The main function, which does initial
228  * functions.
229  */
```

DEBUG CONSOLE PROBLEMS OUTPUT TERMINAL

```
WARNING: Too many lcores enabled. Only 1 used.

Core 0 forwarding packets. [Ctrl+C to quit]
Sending packet: 1
Sending packet: 1
Sending packet: 1
Sending packet: 1
Sending packet: 1
Sending packet: 1
Sending packet: 1
Sending packet: 1
```

识别

- 通过ifconfig找到网卡所对应的ip
- 根据虚拟机的配置找到网卡所对应的网络名: VMnet1
- 运行发包程序, 并在host的wireshark上观察拦截到的包。
- 确认包的各部分信息完整正确。

正在捕获 VMware Network Adapter VMnet1

文件(F) 编辑(E) 视图(V) 跳转(G) 捕获(C) 分析(A) 统计(S) 电话(V) 无线(W) 工具(T) 帮助(H)

应用显示过滤器 -- (Ctrl-F)

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|-----------|---------------|-----------------|----------|--------|---------------------|
| 88 | 35.997534 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 87 | 34.996867 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 86 | 33.995654 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 84 | 32.994792 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 83 | 31.993427 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 82 | 30.991974 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 80 | 29.990774 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 79 | 28.989810 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 78 | 27.988780 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 77 | 26.987643 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 76 | 25.986449 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 75 | 24.985904 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 74 | 23.985326 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 73 | 22.984955 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 72 | 21.984721 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 70 | 20.983799 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 68 | 19.983177 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 66 | 18.982447 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 51 | 17.981933 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 37 | 16.981256 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 36 | 15.980982 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 32 | 14.980323 | 192.168.152.1 | 192.168.0.101 | UDP | 63 | 233 → 233 Len=21 |
| 89 | 36.534995 | 192.168.152.1 | 239.255.255.250 | SSDP | 179 | M-SEARCH * HTTP/1.1 |
| 85 | 33.534547 | 192.168.152.1 | 239.255.255.250 | SSDP | 179 | M-SEARCH * HTTP/1.1 |
| 81 | 30.533999 | 192.168.152.1 | 239.255.255.250 | SSDP | 179 | M-SEARCH * HTTP/1.1 |

> Frame 77: 63 bytes on wire (504 bits), 63 bytes captured (504 bits) on interface \Device\NPF_{4BCDEB48-7585-4FBD-8FBC-49936A8CA880}, id 0
 > Ethernet II, Src: VMware_f2:6b:12 (00:0c:29:f2:6b:12), Dst: VMware_f2:6b:12 (00:0c:29:f2:6b:12)
 > Internet Protocol Version 4, Src: 192.168.152.1, Dst: 192.168.0.101
 > User Datagram Protocol, Src Port: 233, Dst Port: 233
 > Data (21 bytes)

```

0000  00 0c 29 f2 6b 12 00 0c 29 f2 6b 12 00 00 45 00  ..).k...).k...E.
0010  00 31 00 0f 00 00 14 11 0c f6 c0 a8 98 01 c0 a8  .1.....
0020  00 65 00 e9 00 e9 00 1d ff ec 68 65 6c 6c 6f 20  .e.....hello
  
```

VMware Network Adapter VMnet1: (live capture in progress) 分组: 142 · 已显示: 142 (100.0%) 配置: Default